

Task 2: Data Cleansing and Transformation

Instructions:

1. Acquire a real-world dataset requiring data cleaning and transformation.
2. Address data quality issues (missing values, inconsistent formats, outliers).
3. Develop a cleaning strategy (imputation, outlier detection, normalization).
4. Implement necessary transformation steps (feature engineering, aggregation).
5. Validate the cleaned and transformed dataset for integrity and usability.
6. Document the steps taken and provide clear explanations.
7. Present the cleaned and transformed dataset for further analysis.

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('ggplot')
```

```
In [2]: data = pd.read_csv("Datasets/california_housing/housing.csv")
```

1. Acquired Real-world dataset

Dataset details

- Description: The dataset is California Housing Prices dataset.
- Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms, population, households, median_income, median_house_value, ocean_proximity]

```
In [3]: # data.head()
data.info()
```

```
# data['longitude'].isnull().value_counts()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude             20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age    20640 non-null  float64
3   total_rooms           20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population            20640 non-null  float64
6   households            20640 non-null  float64
7   median_income         20640 non-null  float64
8   median_house_value    20640 non-null  float64
9   ocean_proximity       20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

- Only one categorical data Ocean_Proximity

```
In [4]: data['ocean_proximity'].value_counts()
```

```
Out[4]: <1H OCEAN      9136
INLAND          6551
NEAR OCEAN      2658
NEAR BAY        2290
ISLAND           5
Name: ocean_proximity, dtype: int64
```

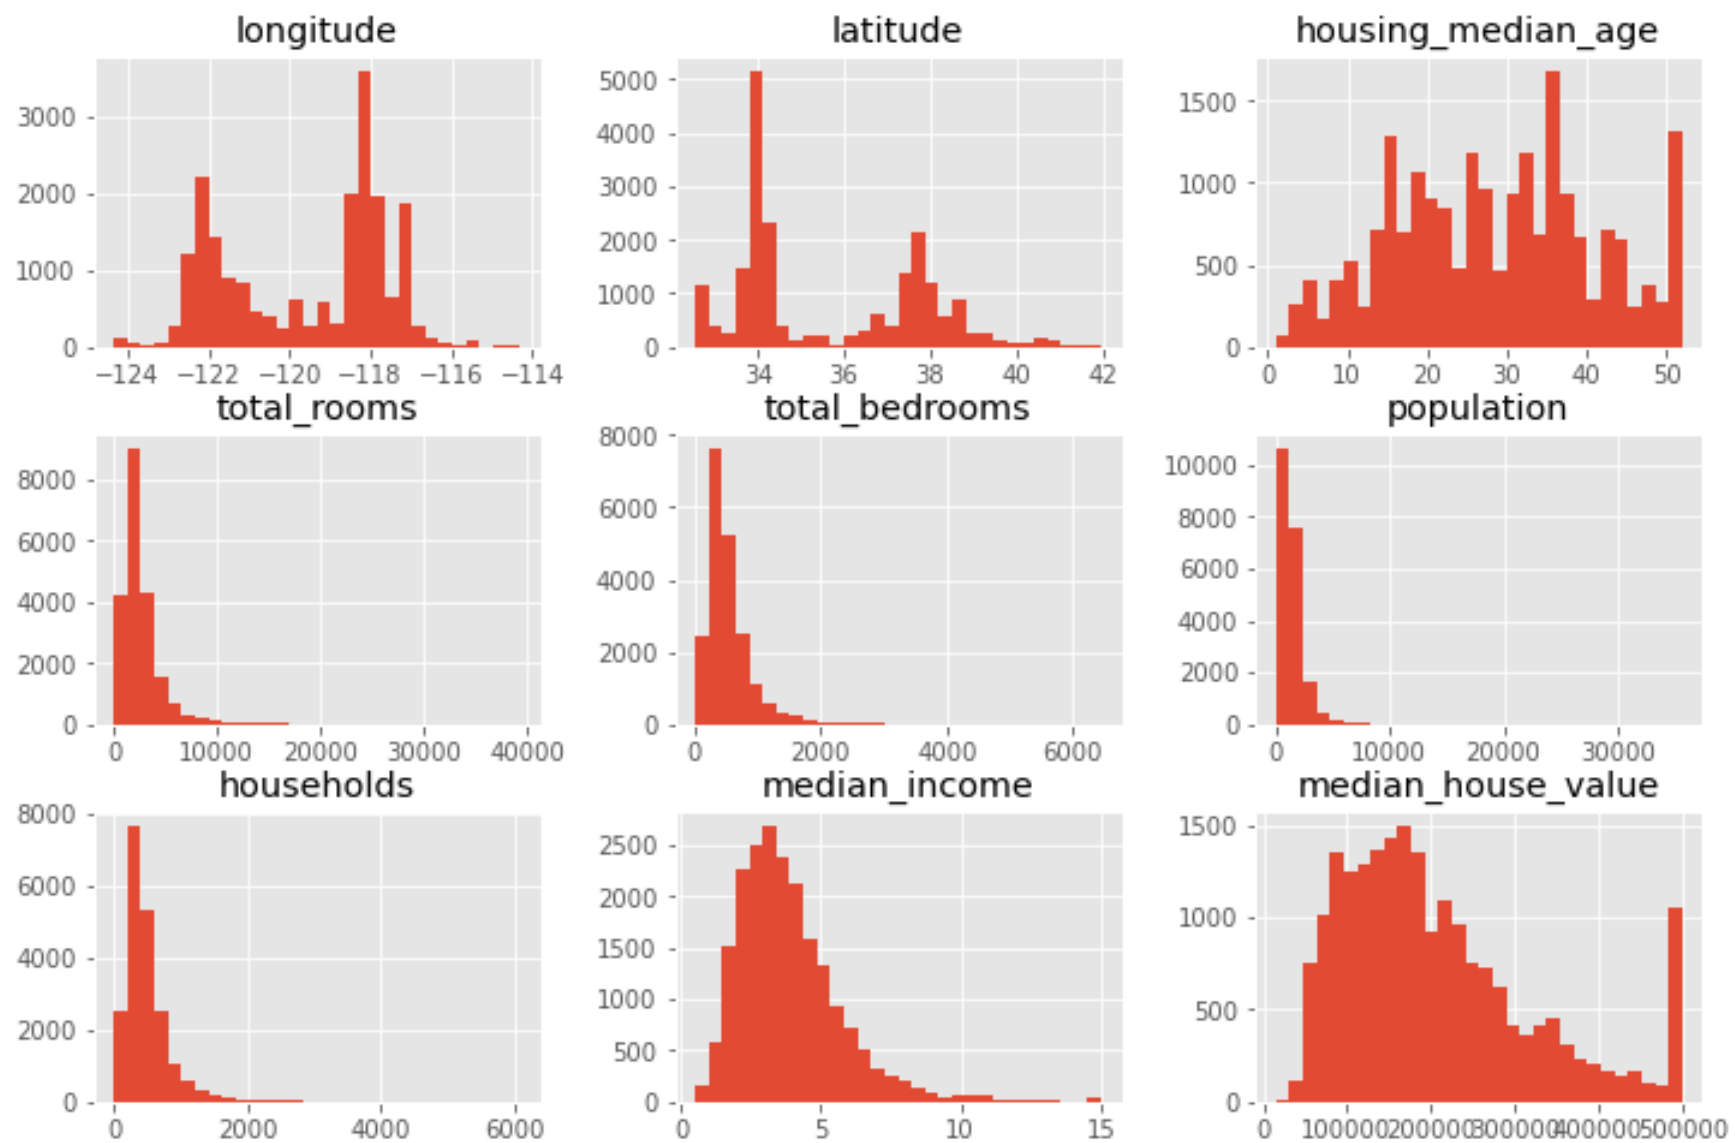
```
In [5]: data.describe()
```

Out[5]:		longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	n
	count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	
	mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	
	std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	
	min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	
	25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	
	50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	
	75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	
	max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	



In [6]: `data.hist(bins=30,figsize=(12,8))`

Out[6]: `array([[<AxesSubplot:title={'center':'longitude'}>,
<AxesSubplot:title={'center':'latitude'}>,
<AxesSubplot:title={'center':'housing_median_age'}>],
[<AxesSubplot:title={'center':'total_rooms'}>,
<AxesSubplot:title={'center':'total_bedrooms'}>,
<AxesSubplot:title={'center':'population'}>],
[<AxesSubplot:title={'center':'households'}>,
<AxesSubplot:title={'center':'median_income'}>,
<AxesSubplot:title={'center':'median_house_value'}>]],
dtype=object)`



2. Data quality and 3. Data cleaning

Pre-processing

- For this task, the data is not normally distributed. Using Stratified sampling technique to prepare the test dataset.
- Creating a new feature `income_label` which is income category and used for sampling.

```
In [7]: data['income_label']=np.ceil(data['median_income']/1.5)
data['income_label'].where(data['income_label']<5,5.0,inplace=True)
```

```
In [8]: from sklearn.model_selection import StratifiedShuffleSplit

split = StratifiedShuffleSplit(n_splits=1,test_size=0.2,random_state=42)

for train_index,test_index in split.split(data,data['income_label']):
    strat_train_set=data.loc[train_index]
    strat_test_set=data.loc[test_index]
```

```
In [9]: strat_train_set.drop('income_label',axis=1,inplace=True)
strat_test_set.drop('income_label',axis=1,inplace=True)
strat_train_set.to_csv("Datasets/california_housing/strat_train_set.csv",index=False)
strat_test_set.to_csv("Datasets/california_housing/strat_test_set.csv",index=False)
```

```
In [24]: data=pd.read_csv('Datasets/california_housing/strat_train_set.csv')
# data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16512 entries, 0 to 16511
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              16512 non-null  float64
1   latitude               16512 non-null  float64
2   housing_median_age     16512 non-null  float64
3   total_rooms            16512 non-null  float64
4   total_bedrooms         16354 non-null  float64
5   population             16512 non-null  float64
6   households             16512 non-null  float64
7   median_income          16512 non-null  float64
8   median_house_value     16512 non-null  float64
9   ocean_proximity        16512 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.3+ MB
```

EDA

- Carrying out various visualization on train dataset for realising patterns, correlations and getting the sense of the data

```
In [11]: plt.figure(figsize=(10,6))
plt.scatter(x=data['longitude'],y=data['latitude'])
plt.title("Distribution of households",size=16)
```

```
Out[11]: Text(0.5, 1.0, 'Distribution of households')
```

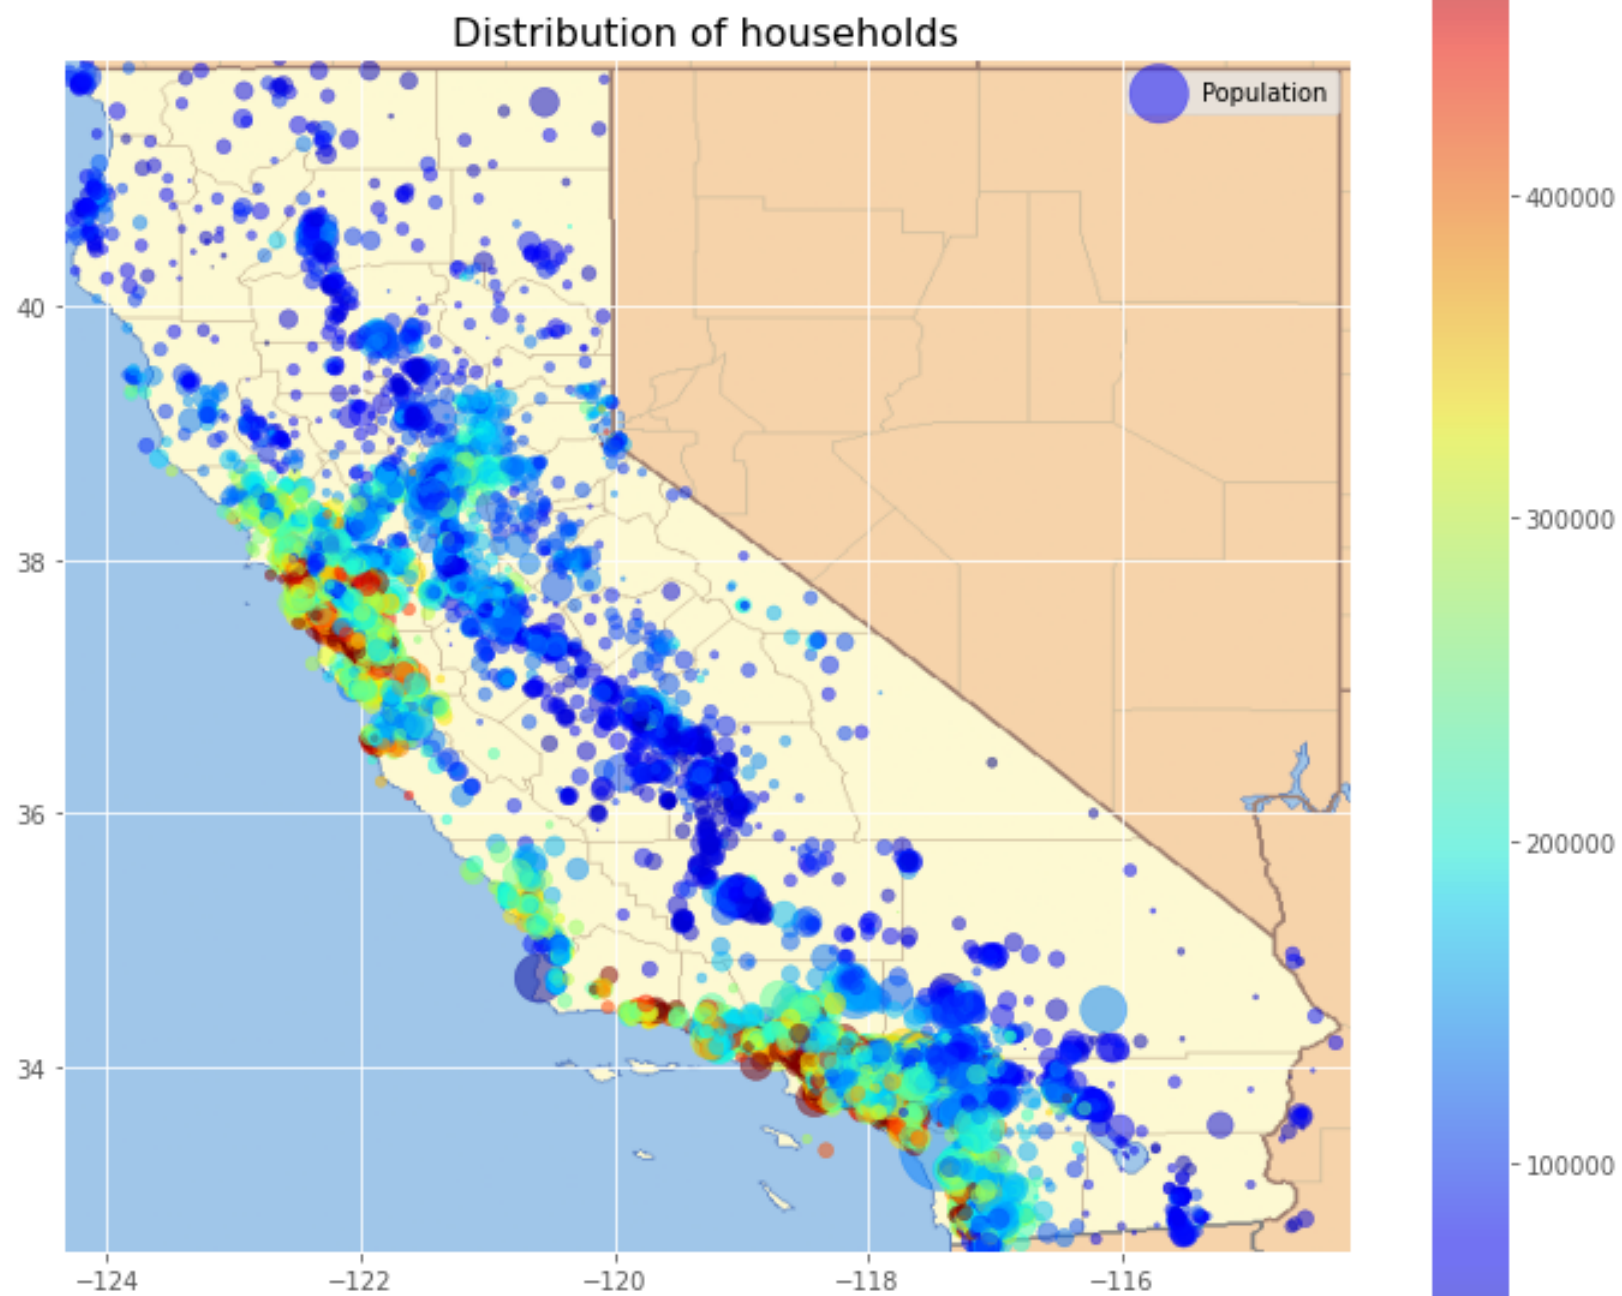


```
In [12]: plt.figure(figsize=(12,12))
img=plt.imread('Datasets/california_housing/california.png')
plt.imshow(img,zorder=0,extent=[-124.35,-114.2,32.54,41.95])

plt.scatter(x=data['longitude'],y=data['latitude'],alpha=0.5,s=data['population']/30,c=data['median_house_value'])
plt.colorbar()
plt.title("Distribution of households",size=16)
plt.legend()
```

```
C:\Users\MAHAVIR\AppData\Local\Temp\ipykernel_5504\1422664156.py:6: MatplotlibDeprecationWarning: Auto-removal of  
grids by pcolor() and pcolormesh() is deprecated since 3.5 and will be removed two minor releases later; please c  
all grid(False) first.  
    plt.colorbar()
```

```
Out[12]: <matplotlib.legend.Legend at 0x1fad86b2c70>
```



Inights from Visualization

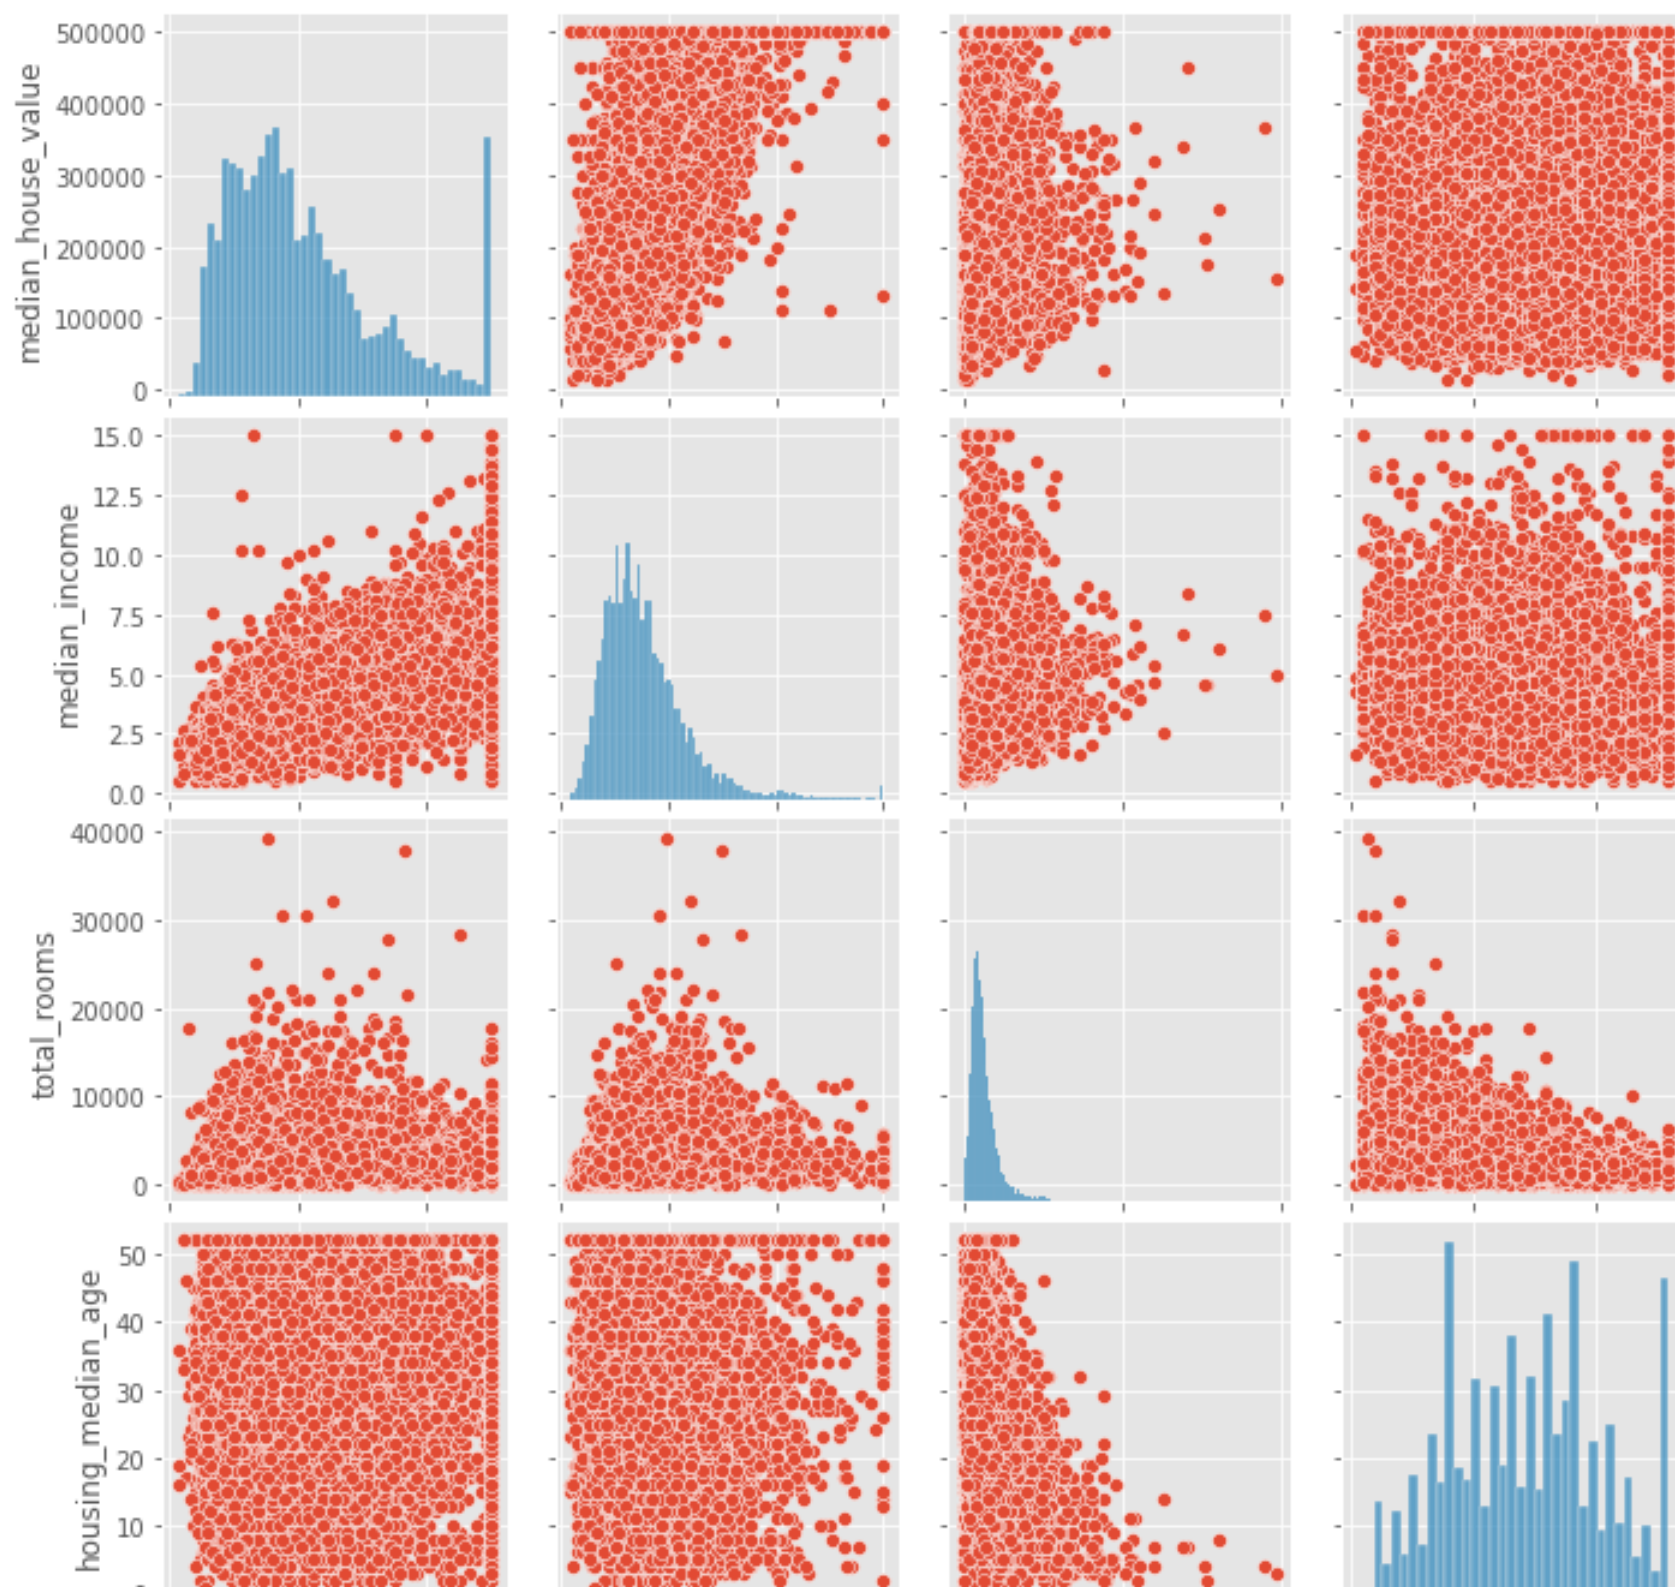
- Housing prices are much related to location and population density.
- Housing prices near ocean are higher except in northern california.
- Now, see the correlation of 'medial house value' with other columns. This is Pearson's correlation coefficient.

```
In [13]: corr_matrix=data.corr()  
corr_matrix['median_house_value'].sort_values(ascending=False)
```

```
Out[13]: median_house_value    1.000000  
median_income      0.687151  
total_rooms        0.135140  
housing_median_age  0.114146  
households         0.064590  
total_bedrooms     0.047781  
population        -0.026882  
longitude          -0.047466  
latitude           -0.142673  
Name: median_house_value, dtype: float64
```

```
In [14]: sns.pairplot(data[['median_house_value', 'median_income', 'total_rooms', 'housing_median_age']])
```

```
Out[14]: <seaborn.axisgrid.PairGrid at 0x1fad7374400>
```






- Median Income is the most promising attribute to get Median Hosung Price

```
In [15]: plt.figure(figsize=(10,6))  
plt.scatter(y=data['median_house_value'],x=data['median_income'],alpha=0.3)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x1fadb55e8e0>
```



- A clear line can be seen at 500k at which the data is capped. Similar lines can be seen around 450k,350k. This kind of data may degrade the performance of model.

4. Feature Engineering

Creating new features:

- rooms per household
- bedrooms per room
- population per household

```
In [16]: data_copy = data.copy()
```

```
In [17]: data_copy['rooms_per_household']=data_copy['total_rooms']/data_copy['households']
data_copy['bedrooms_per_room']=data_copy['total_bedrooms']/data_copy['total_rooms']
data_copy['population_per_household']=data_copy['population']/data_copy['households']
# data_copy.head()
```

```
In [18]: corr_matrix=data_copy.corr()
corr_matrix['median_house_value'].sort_values(ascending=False)
```

```
Out[18]: median_house_value      1.000000
median_income      0.687151
rooms_per_household  0.146255
total_rooms      0.135140
housing_median_age  0.114146
households      0.064590
total_bedrooms      0.047781
population_per_household -0.021991
population      -0.026882
longitude      -0.047466
latitude      -0.142673
bedrooms_per_room  -0.259952
Name: median_house_value, dtype: float64
```

5. Validation of new features and data


- It is clear that rooms_per_household and bedrooms_per_room have better correlation with median_house_value than total_rooms and total_bedrooms.
- Later Feature adder class needs to be created later during testing.(transforming the data to have new features)


```
In [19]: data.isnull().value_counts()
```

```
Out[19]: longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  households  median_income  medi
an_house_value  ocean_proximity
False          False    False                False        False          False        False        False        Fals
e              False                16354                                True          False        False        False        Fals
e              False                158
dtype: int64
```

```
In [20]: data_copy.hist(column='total_bedrooms',bins=30)
data_copy.hist(column='bedrooms_per_room',bins=30)
```

```
Out[20]: array([[<AxesSubplot:title={'center':'bedrooms_per_room'}>]], dtype=object)
```




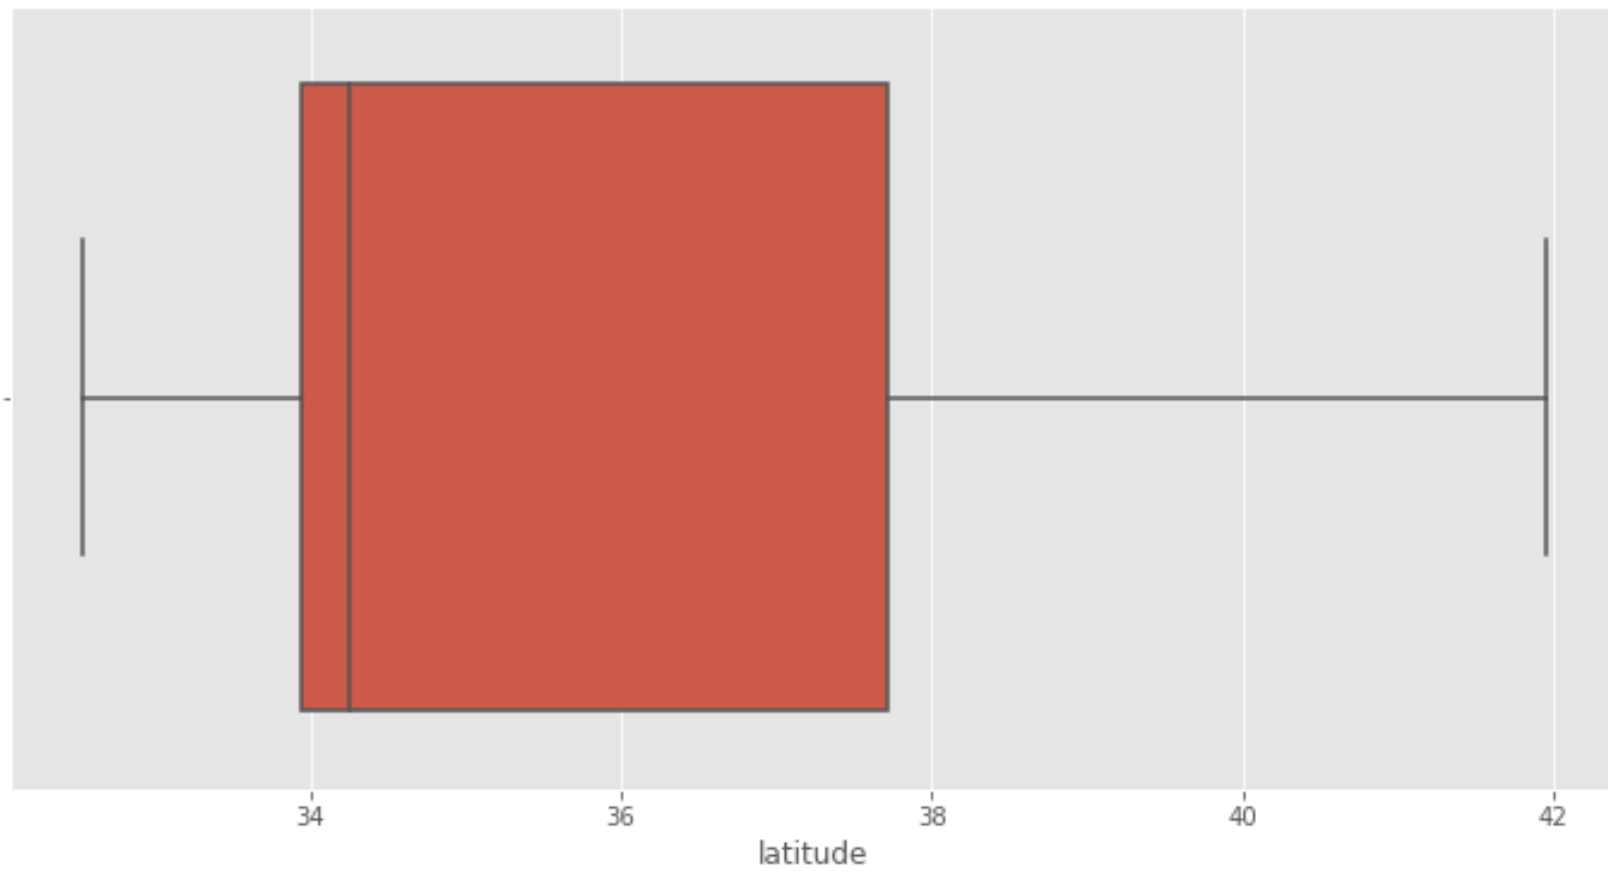



Cleaning the data


- Removing Outliers realised through boxplots

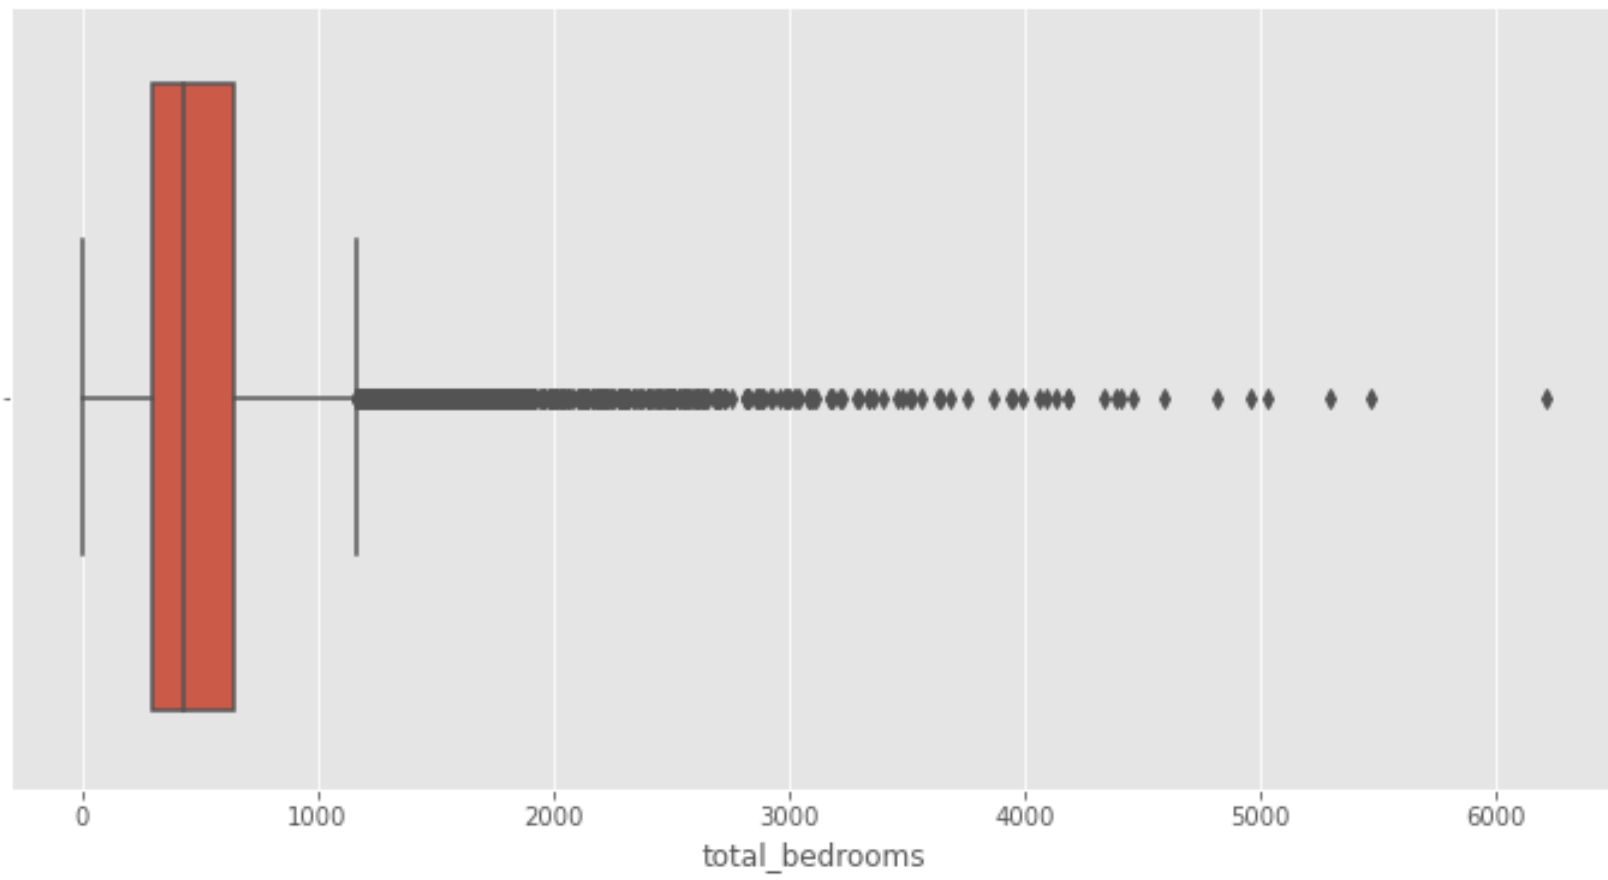
```
In [21]: num_features=['longitude', 'latitude', 'housing_median_age', 'total_rooms',  
                    'total_bedrooms', 'population', 'households', 'median_income',  
                    'median_house_value', 'rooms_per_household',  
                    'bedrooms_per_room', 'population_per_household']  
  
for i in num_features:  
    fig, ax = plt.subplots()  
    fig.set_size_inches(12,6)  
    #plt.xlim(-10,10)  
    sns.boxplot(x=i,data=data_copy,ax=ax)
```










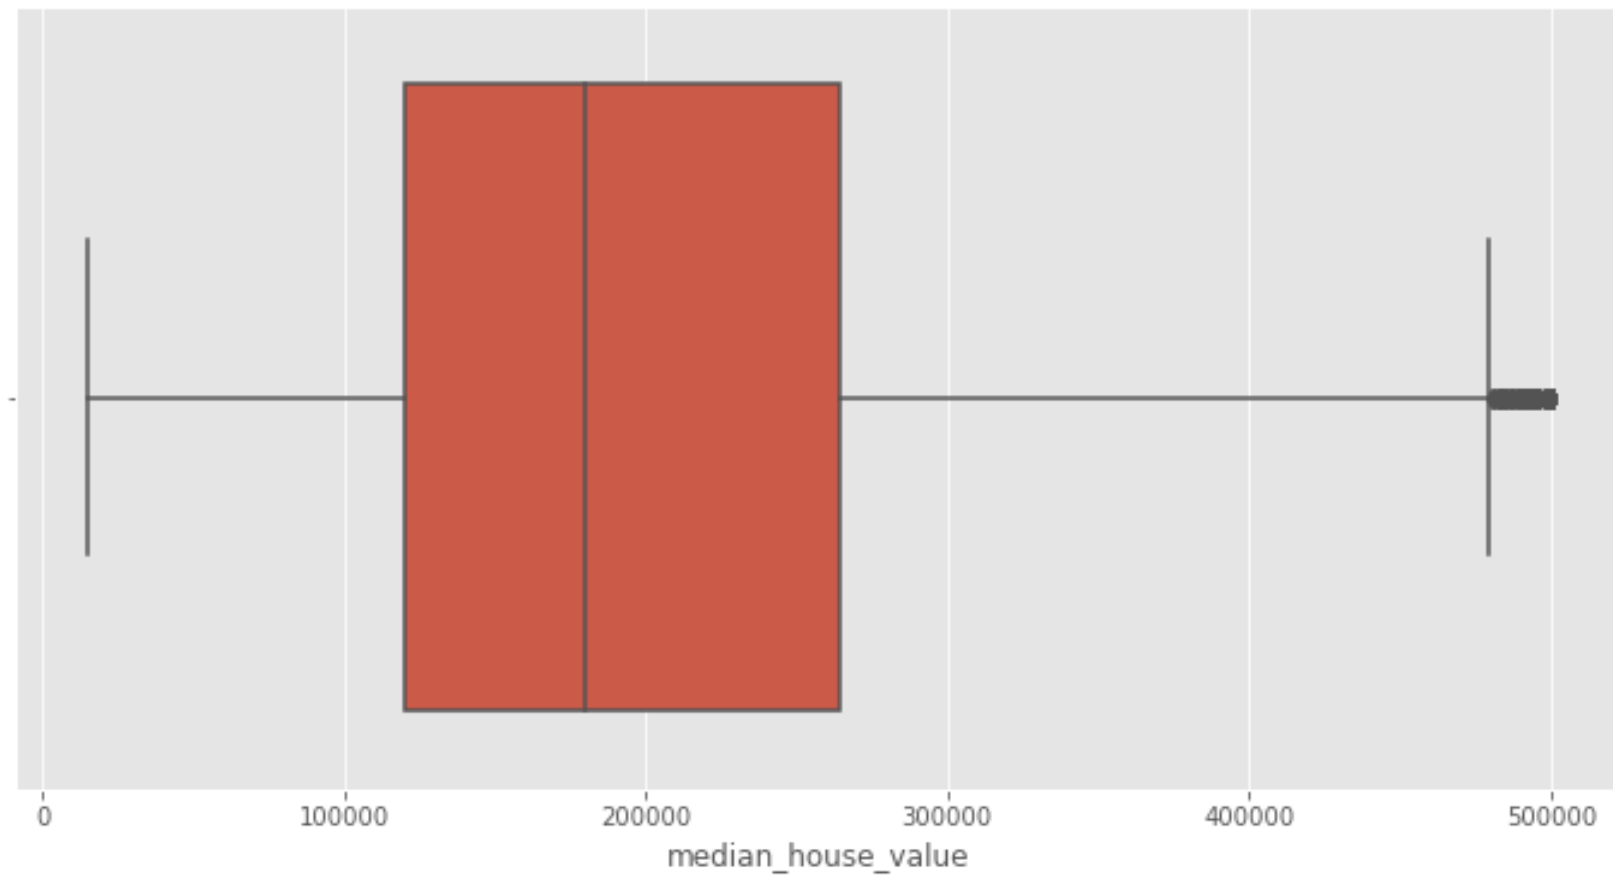





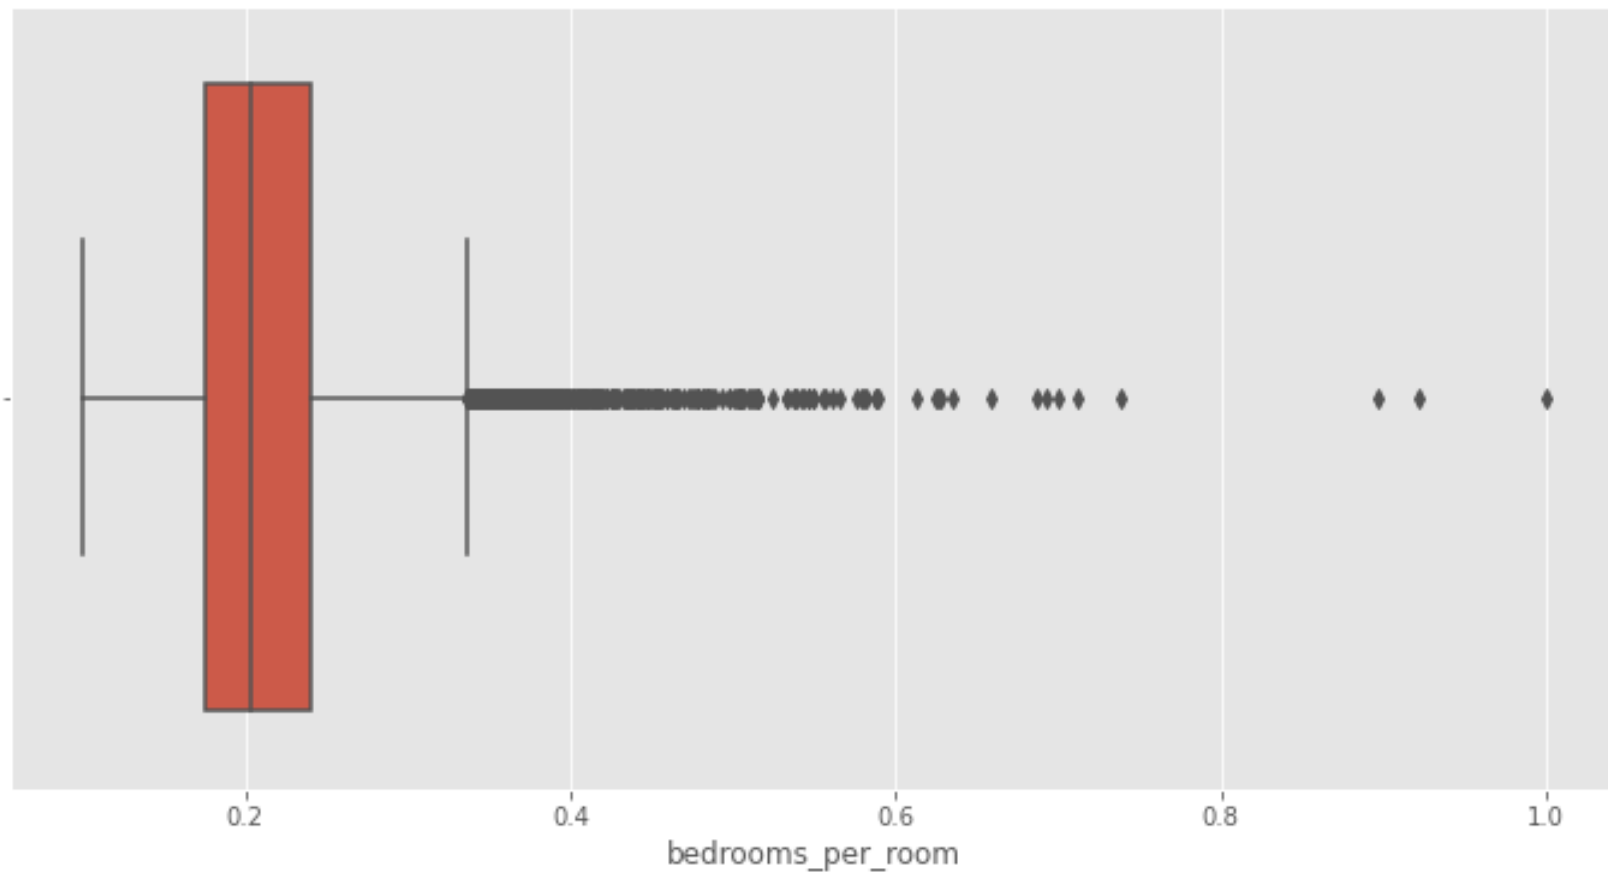


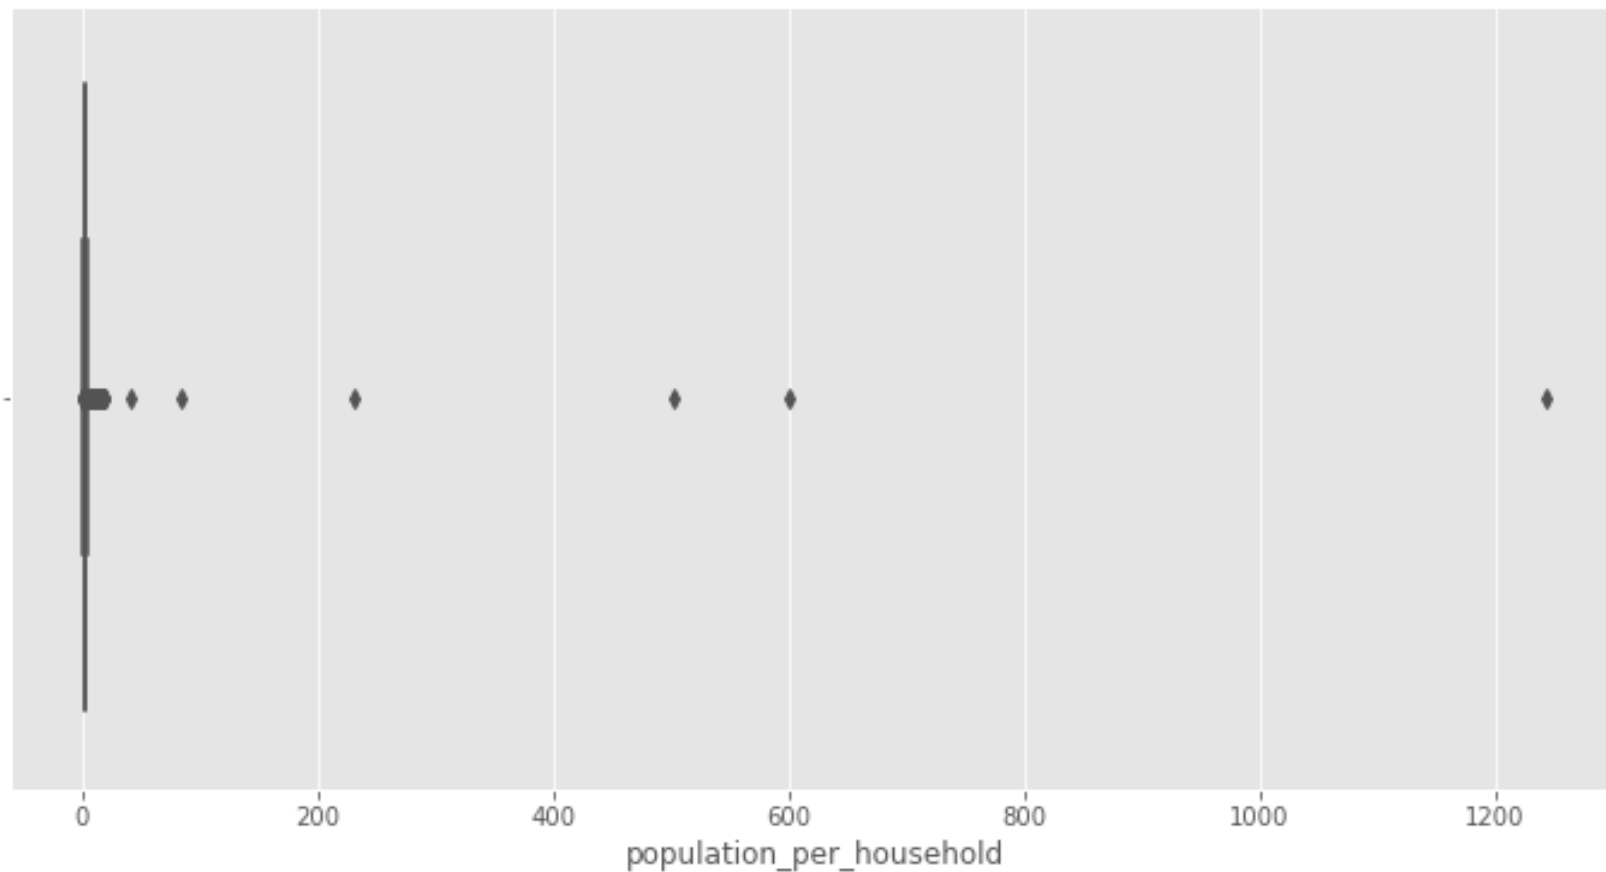












In [22]: `from sklearn.base import BaseEstimator,TransformerMixin`

```
class RemoveOutliers(BaseEstimator,TransformerMixin):
    """This class removes outliers from data.
    Note: Outlier values are hard coded
    """
    def fit (self,X,y=None):
        return self

    def transform(self,X,y=None):
        X=X[(X['median_house_value']!=500001) | (X['median_income']>=2)].reset_index(drop=True)
        X=X[X['median_income']<=11].reset_index(drop=True)
        X=X[(X['median_house_value']!=350000) | (X['median_income']>=1.5)].reset_index(drop=True)
        X=X[(X['median_house_value']!=450000) | (X['median_income']>=2)].reset_index(drop=True)
        X=X[(X['median_house_value']>=350000) | (X['median_income']<=9.5)].reset_index(drop=True)
        X=X[X['population']<=9000]
```

```

X=X[(X['population_per_household']>=1.15) & (X['population_per_household']<=6.5)]
X=X[X['rooms_per_household']<20]
X=X[X['bedrooms_per_room']<0.5].reset_index(drop=True)
return X

```

```
data_copy=RemoveOutliers().fit_transform(data_copy)
```

```

data_labels=data_copy['median_house_value']
data_copy=data_copy.drop('median_house_value',axis=1)

```

```

In [23]: # data_copy.head()
# data_copy.isnull().value_counts()
data_copy.info()
data_copy.hist(bins=50,figsize=(15,12))












```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16005 entries, 0 to 16004
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   longitude             16005 non-null  float64
 1   latitude              16005 non-null  float64
 2   housing_median_age    16005 non-null  float64
 3   total_rooms           16005 non-null  float64
 4   total_bedrooms        16005 non-null  float64
 5   population             16005 non-null  float64
 6   households            16005 non-null  float64
 7   median_income         16005 non-null  float64
 8   ocean_proximity       16005 non-null  object
 9   rooms_per_household   16005 non-null  float64
10   bedrooms_per_room     16005 non-null  float64
11   population_per_household 16005 non-null  float64
dtypes: float64(11), object(1)
memory usage: 1.5+ MB

```

```
Out[23]: array([[<AxesSubplot:title={'center':'longitude'}>,
<AxesSubplot:title={'center':'latitude'}>,
<AxesSubplot:title={'center':'housing_median_age'}>],
[<AxesSubplot:title={'center':'total_rooms'}>,
<AxesSubplot:title={'center':'total_bedrooms'}>,
<AxesSubplot:title={'center':'population'}>],
[<AxesSubplot:title={'center':'households'}>,
<AxesSubplot:title={'center':'median_income'}>,
<AxesSubplot:title={'center':'rooms_per_household'}>],
[<AxesSubplot:title={'center':'bedrooms_per_room'}>,
<AxesSubplot:title={'center':'population_per_household'}>,
<AxesSubplot:>]], dtype=object)
```



6. Summary of steps taken

Steps undertaken for pre-processing the data:

- Data has been cleaned with no null values and outliers.
- We have further converted the categorical feature to numeric and scaled the data. ##### Further Analysis steps to be done:
- Remove skewness
- Can also use Get dummies to convert categorical feature of 'ocean_proximity'
- Check for Multi colinearity and scale the features further.